

Shrinkage, False Discovery Rates, and an Alternative to the Zero Assumption

Matthew Stephens

2013/11/1

The Canonical Genomics Experiment

- Measure lots of things, with error

The Canonical Genomics Experiment

- Measure lots of things, with error
- Get estimates of effects β_j ($\hat{\beta}_j$) and their standard errors s_j

The Canonical Genomics Experiment

- Measure lots of things, with error
- Get estimates of effects β_j ($\hat{\beta}_j$) and their standard errors s_j
- Turn these into Z-scores, $z_j = \hat{\beta}_j/s_j$

The Canonical Genomics Experiment

- Measure lots of things, with error
- Get estimates of effects β_j ($\hat{\beta}_j$) and their standard errors s_j
- Turn these into Z-scores, $z_j = \hat{\beta}_j/s_j$
- Turn these into p values, p_j

The Canonical Genomics Experiment

- Measure lots of things, with error
- Get estimates of effects β_j ($\hat{\beta}_j$) and their standard errors s_j
- Turn these into Z-scores, $z_j = \hat{\beta}_j/s_j$
- Turn these into p values, p_j
- Apply `qvalue` to identify findings “significant” at a given False Discovery Rate.

The Canonical Genomics Experiment

- Measure lots of things, with error
- Get estimates of effects β_j ($\hat{\beta}_j$) and their standard errors s_j
- Turn these into Z-scores, $z_j = \hat{\beta}_j/s_j$
- Turn these into p values, p_j
- Apply `qvalue` to identify findings “significant” at a given False Discovery Rate.
- ...?

FDR, local fdr, and q values

Although precise definitions vary depending on whether one takes a Bayesian or Frequentist approach to the problem, roughly

- The FDR at a threshold P is

$$\text{FDR}(P) = \Pr(\beta_j = 0 | p_j < P).$$

FDR, local fdr, and q values

Although precise definitions vary depending on whether one takes a Bayesian or Frequentist approach to the problem, roughly

- The FDR at a threshold P is

$$\text{FDR}(P) = \Pr(\beta_j = 0 | p_j < P).$$

- The q value for observation j is $q_j = \text{FDR}(p_j)$.

FDR, local fdr, and q values

Although precise definitions vary depending on whether one takes a Bayesian or Frequentist approach to the problem, roughly

- The FDR at a threshold P is

$$\text{FDR}(P) = \Pr(\beta_j = 0 | p_j < P).$$

- The q value for observation j is $q_j = \text{FDR}(p_j)$.
- The local false discovery rate, fdr , at threshold P is

$$\text{fdr}(P) = \Pr(\beta_j = 0 | p_j = P).$$

FDR, local fdr, and q values

Although precise definitions vary depending on whether one takes a Bayesian or Frequentist approach to the problem, roughly

- The FDR at a threshold P is


$$\text{FDR}(P) = \Pr(\beta_j = 0 | p_j < P).$$

- The q value for observation j is $q_j = \text{FDR}(p_j)$.
- The local false discovery rate, fdr , at threshold P is

$$\text{fdr}(P) = \Pr(\beta_j = 0 | p_j = P).$$


- The fdr is more relevant, but slightly harder to estimate than FDR because it involves density estimation rather than tail-area estimation.

Example: FDR estimation




Example: FDR estimation

Distribution of p values




Data from Hedenfalk et al. comparing BRCA1 vs BRCA2 expression


Example: FDR estimation




Example: fdr estimation



Example: fdr estimation



Example: fdr estimation



FDR problem 1: different measurement precision

- If some effects are measured very imprecisely, those tests “lack power” and simply add noise

FDR problem 1: different measurement precision

- If some effects are measured very imprecisely, those tests “lack power” and simply add noise
- In particular, such tests increase the estimated number of nulls, and increase the FDR for other tests

FDR problem 1: different measurement precision


- If some effects are measured very imprecisely, those tests “lack power” and simply add noise
- In particular, such tests increase the estimated number of nulls, and increase the FDR for other tests
- It would seem preferable to simply ignore the tests with very low precision. Summarizing each test by a p value (or Z score) loses the information about precision.

Example: Mouse Heart Data


##	gene	lv1	lv2	rv1	rv2	genelength
## 1	Itm2a	2236	2174	9484	10883	1626
## 2	Sergef	97	90	341	408	1449
## 3	Fam109a	383	314	1864	2384	2331
## 4	Dhx9	2688	2631	18501	20879	4585
## 5	Ssu72	762	674	2806	3435	1446
## 8	Eif2b2	736	762	3081	3601	1565

- Data on 150 mouse hearts, dissected into left and right ventricle (courtesy Scott Schmemo, Marcelo Nobrega)


Example: Mouse Heart Data




Mouse Data: Counts vary considerably across genes




Lower count genes, less power



Higher count genes, more power



FDR problem 1: low count genes add noise, increase q values



FDR problem 1: Summary

- Analyzing p values or Z scores doesn't fully account for measurement precision.

Problem 2: The Zero Assumption (ZA)

- The standard q value approach assumes that all the p values near 1 are null.

Problem 2: The Zero Assumption (ZA)

- The standard q value approach assumes that all the p values near 1 are null.
- Analogously, one can assume that all Z scores near 0 are null. Efron refers to this as the “Zero Assumption”.

Problem 2: The Zero Assumption (ZA)

- The standard q value approach assumes that all the p values near 1 are null.
- Analogously, one can assume that all Z scores near 0 are null. Efron refers to this as the “Zero Assumption”.
- The ZA allows us to estimate the null proportion, π_0 , using the density of p values near 1 (or Z scores near 0).


Problem 2: The ZA

- The ZA seems initially natural.


Problem 2: The ZA

- The ZA seems initially natural.
- However, it turns out to imply unrealistic assumptions about the distribution of non-zero effects.


Implied distribution of p values under H_1




Implied distribution of Z scores under alternative (fdrtool)



Implied distribution of Z scores under alternative (locfdr)



Implied distribution of Z scores under alternative (mixfdr)



Problems: Summary

- By summarizing each observation by a Z score or p value, standard fdr tools ignore precision of different measurements

Problems: Summary

- By summarizing each observation by a Z score or p value, standard fdr tools ignore precision of different measurements
- Standard tools make the ZA, which implies actual effects have a (probably unrealistic) bimodal distribution. [and tends to overestimate π_0 , losing power]

Problems: Summary

- By summarizing each observation by a Z score or p value, standard fdr tools ignore precision of different measurements
- Standard tools make the ZA, which implies actual effects have a (probably unrealistic) bimodal distribution. [and tends to overestimate π_0 , losing power]
- Also standard tools focus only on zero vs non-zero effects. (eg what if we would like to identify genes that have at least a 2-fold change?)

FDR via Empirical Bayes

- Following previous work (e.g. Newton, Efron, Muralidharan) we take an empirical Bayes approach to FDR.

FDR via Empirical Bayes

- Following previous work (e.g. Newton, Efron, Muralidharan) we take an empirical Bayes approach to FDR.
- Eg Efron assumes that the Z scores come from a mixture of null, and alternative:

$$Z_j \sim f_Z(.) = \pi_0 N(., 0, 1) + (1 - \pi_0) f_1(.)$$

where f_1 is to be estimated from the data.

FDR via Empirical Bayes

- Following previous work (e.g. Newton, Efron, Muralidharan) we take an empirical Bayes approach to FDR.
- Eg Efron assumes that the Z scores come from a mixture of null, and alternative:

$$Z_j \sim f_Z(.) = \pi_0 N(., 0, 1) + (1 - \pi_0) f_1(.)$$

where f_1 is to be estimated from the data.

- Various semi-parametric approaches taken to estimating f_1 . For example, Efron uses Poisson regression; Muralidharan uses mixture of normal distributions.

FDR via Empirical Bayes

- Following previous work (e.g. Newton, Efron, Muralidharan) we take an empirical Bayes approach to FDR.
- Eg Efron assumes that the Z scores come from a mixture of null, and alternative:

$$Z_j \sim f_Z(.) = \pi_0 N(., 0, 1) + (1 - \pi_0) f_1(.)$$

where f_1 is to be estimated from the data.

- Various semi-parametric approaches taken to estimating f_1 . For example, Efron uses Poisson regression; Muralidharan uses mixture of normal distributions.
- $\text{fdr}(Z) \approx \pi_0 N(Z; 0, 1) / f_Z(Z)$

FDR: The New Deal

- Instead of modelling Z scores, model the effects β ,

$$\beta_j \sim \pi_0 \delta_0(.) + (1 - \pi_0)g(.)$$

FDR: The New Deal

- Instead of modelling Z scores, model the effects β ,

$$\beta_j \sim \pi_0 \delta_0(.) + (1 - \pi_0)g(.)$$

- Constrain g to be unimodal about 0.

FDR: The New Deal

- Instead of modelling Z scores, model the effects β ,

$$\beta_j \sim \pi_0 \delta_0(.) + (1 - \pi_0)g(.)$$

- Constrain g to be unimodal about 0.
- *Incorporate precision* of each observation $\hat{\beta}$ into the likelihood. Specifically, approximate likelihood for β_j by a normal:

$$L(\beta_j) \propto \exp(-0.5(\beta_j - \hat{\beta}_j)^2/s_j^2).$$

[From $\hat{\beta}_j \sim N(\beta_j, s_j)$]

FDR: The New Deal

- Instead of modelling Z scores, model the effects β ,

$$\beta_j \sim \pi_0 \delta_0(.) + (1 - \pi_0)g(.)$$

- Constrain g to be unimodal about 0.
- *Incorporate precision* of each observation $\hat{\beta}$ into the likelihood. Specifically, approximate likelihood for β_j by a normal:

$$L(\beta_j) \propto \exp(-0.5(\beta_j - \hat{\beta}_j)^2/s_j^2).$$

[From $\hat{\beta}_j \sim N(\beta_j, s_j)$]

- fdr given by

$$p(\beta_j = 0|\hat{\beta}_j) = \pi_0 p(\hat{\beta}_j|\beta_j = 0)/p(\hat{\beta}_j)$$

FDR - The New Deal

- A convenient way to model g is by a mixture of 0-centered normal distributions:

$$g(\beta; \pi) = \sum_{k=1}^K \pi_k N(\beta; 0, \sigma_k^2)$$

FDR - The New Deal

- A convenient way to model g is by a mixture of 0-centered normal distributions:

$$g(\beta; \pi) = \sum_{k=1}^K \pi_k N(\beta; 0, \sigma_k^2)$$

- Estimating g comes down to estimating π . Joint estimation of π_0, π easy by maximum likelihood (EM algorithm) or variational Bayes.

FDR - The New Deal

- A convenient way to model g is by a mixture of 0-centered normal distributions:

$$g(\beta; \pi) = \sum_{k=1}^K \pi_k N(\beta; 0, \sigma_k^2)$$

- Estimating g comes down to estimating π . Joint estimation of π_0, π easy by maximum likelihood (EM algorithm) or variational Bayes.
- By allowing K large, and σ_k to span a dense grid of values, we get a fairly flexible unimodal symmetric distribution.

FDR - The New Deal

- A convenient way to model g is by a mixture of 0-centered normal distributions:

$$g(\beta; \pi) = \sum_{k=1}^K \pi_k N(\beta; 0, \sigma_k^2)$$

- Estimating g comes down to estimating π . Joint estimation of π_0, π easy by maximum likelihood (EM algorithm) or variational Bayes.
- By allowing K large, and σ_k to span a dense grid of values, we get a fairly flexible unimodal symmetric distribution.
- Alternatively, a mixture of uniforms, with 0 as one end-point of the range, provides still more flexibility, and in particular allows for asymmetry.

FDR - The New Deal

- A convenient way to model g is by a mixture of 0-centered normal distributions:

$$g(\beta; \pi) = \sum_{k=1}^K \pi_k N(\beta; 0, \sigma_k^2)$$

- Estimating g comes down to estimating π . Joint estimation of π_0, π easy by maximum likelihood (EM algorithm) or variational Bayes.
- By allowing K large, and σ_k to span a dense grid of values, we get a fairly flexible unimodal symmetric distribution.
- Alternatively, a mixture of uniforms, with 0 as one end-point of the range, provides still more flexibility, and in particular allows for asymmetry.
- If allow a very large number of uniforms this provides the non-parametric mle for g ; cf Grenander 1953; Campy + Thomas.

Illustration: g a mixture of 0-centered normals




Illustration: g a mixture of 0-centered normals




Illustration: g a mixture of 0-anchored uniforms





Illustration: g a mixture of 0-anchored uniforms



Issue: identifiability of π_0

- For estimating False Discoveries, we are asking whether $\beta_j = 0$.

Issue: identifiability of π_0

- For estimating False Discoveries, we are asking whether $\beta_j = 0$.
- However, the data cannot distinguish between $\beta_j = 0$ and β_j “very small”

Issue: identifiability of π_0

- For estimating False Discoveries, we are asking whether $\beta_j = 0$.
- However, the data cannot distinguish between $\beta_j = 0$ and β_j “very small”
- As a result π_0 is formally unidentifiable. Eg data can never rule out $\pi_0 = 0$.

Issue: identifiability of π_0

- The Zero assumption (ZA) solves the identifiability problem by assuming that there *are* no β_j near zero!

Issue: identifiability of π_0

- The Zero assumption (ZA) solves the identifiability problem by assuming that there *are* no β_j near zero!
- The ZA makes π_0 identifiable.

Issue: identifiability of π_0

- The Zero assumption (ZA) solves the identifiability problem by assuming that there *are* no β_j near zero!
- The ZA makes π_0 identifiable.
- Another view is that the estimate of π_0 under ZA will systematically tend to overestimate π_0 , and so is “conservative”.

Issue: identifiability of π_0

- The Zero assumption (ZA) solves the identifiability problem by assuming that there *are* no β_j near zero!
- The ZA makes π_0 identifiable.
- Another view is that the estimate of π_0 under ZA will systematically tend to overestimate π_0 , and so is “conservative”.
- That is it provides an “upper bound” on π_0

Identifiability of π_0 : Solution 1

- We replaced the ZA with the unimodal assumption on g .

Identifiability of π_0 : Solution 1

- We replaced the ZA with the unimodal assumption on g .
- This does not make π_0 identifiable, but it does effectively provide an upper bound on π_0 .

Identifiability of π_0 : Solution 1

- We replaced the ZA with the unimodal assumption on g .
- This does not make π_0 identifiable, but it does effectively provide an upper bound on π_0 .
- Indeed, we saw that when we estimated π_0 under the ZA the data then contradicted the unimodal assumption on g . Thus the upper bound is more conservative than under ZA.


Identifiability of π_0 : Solution 1

- We replaced the ZA with the unimodal assumption on g .
- This does not make π_0 identifiable, but it does effectively provide an upper bound on π_0 .
- Indeed, we saw that when we estimated π_0 under the ZA the data then contradicted the unimodal assumption on g . Thus the upper bound is more conservative than under ZA.
- In practice, implement upper bound by putting prior on π_0 that encourages it to be big, then estimate π by posterior mean (VB).

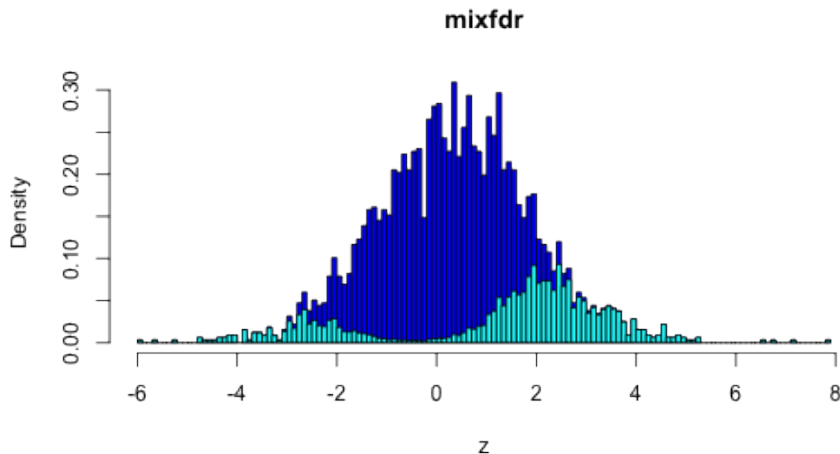
Illustration: BRCA data

Example: BRCA data

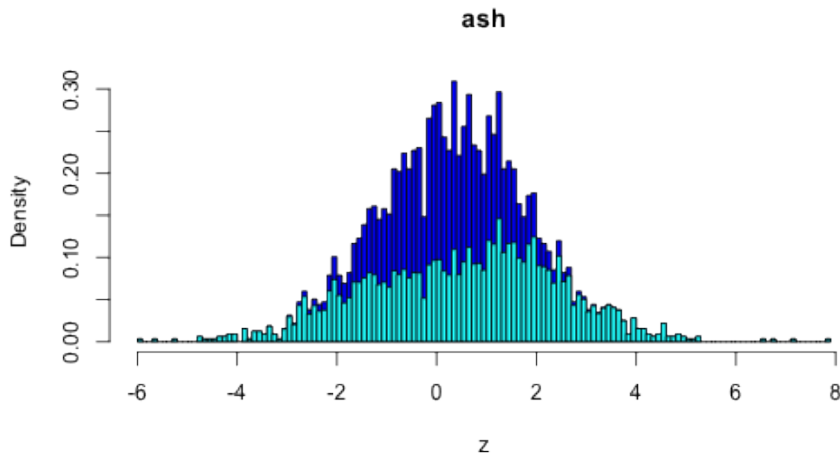
Compare fitted $f(\beta)$, both estimating π_0 and fixing $\pi_0 = 0$.



Recall Problem: distribution of alternative Z values multimodal



Problem Fixed: distribution of alternative Z values unimodal



BRCA1: Compare π_0 estimates

```
round(c(hh.fdrtool$param[3], hh.locfdr$fp0[1, 3], hh.mixfdr$pi0[1, 2])
```

```
## [1] 0.64 0.74 0.80 0.47
```


BRCA1: Compare number significant at $\text{fdr} < 0.05$

```
c(sum(hh.fdrtool$lfd < 0.05), sum(hh.locfdr$fdr < 0.05), sum(hh.fdrtool$ZeroProb < 0.05), sum(hh.ashz$ZeroProb < 0.05))
```

```
## [1] 154 171 162 197
```

Identifiability of π_0 : Solution 2

- Identifiability of π_0 is primarily an issue if we insist on asking question is $\beta_j = 0$?

Identifiability of π_0 : Solution 2

- Identifiability of π_0 is primarily an issue if we insist on asking question is $\beta_j = 0$?
- How about we change focus: assume *none* of the β_j are zero (“one group approach”), and ask for which β_j are we confident about the sign (Gelman et al, 2012).

Identifiability of π_0 : Solution 2

- Identifiability of π_0 is primarily an issue if we insist on asking question is $\beta_j = 0$?
- How about we change focus: assume *none* of the β_j are zero (“one group approach”), and ask for which β_j are we confident about the sign (Gelman et al, 2012).
- Positive and negative effects are often treated differently in practice anyway.

Identifiability of π_0 : Solution 2

- Identifiability of π_0 is primarily an issue if we insist on asking question is $\beta_j = 0$?
- How about we change focus: assume *none* of the β_j are zero (“one group approach”), and ask for which β_j are we confident about the sign (Gelman et al, 2012).
- Positive and negative effects are often treated differently in practice anyway.
- That is we replace fdr with False Sign Rate (fsr), the probability that if we say an effect is positive (negative), it is not.

Identifiability of π_0 : Solution 2

- Identifiability of π_0 is primarily an issue if we insist on asking question is $\beta_j = 0$?
- How about we change focus: assume *none* of the β_j are zero (“one group approach”), and ask for which β_j are we confident about the sign (Gelman et al, 2012).
- Positive and negative effects are often treated differently in practice anyway.
- That is we replace fdr with False Sign Rate (fsr), the probability that if we say an effect is positive (negative), it is not.
- Example: suppose we estimate that $\Pr(\beta_j < 0) = 0.975$ and $\Pr(\beta_j > 0) = 0.025$. Then we report β_j as a “(negative) discovery“, and estimate its fsr as 0.025.

A technicality

- Suppose you estimate $\Pr(\beta_j < 0) = 0.98$, $\Pr(\beta_j > 0) = 0.01$, $\Pr(\beta_j = 0) = 0.01$.

A technicality

- Suppose you estimate $\Pr(\beta_j < 0) = 0.98$, $\Pr(\beta_j > 0) = 0.01$, $\Pr(\beta_j = 0) = 0.01$.
- Should you declare an fdr of 0.01 or 0.02?

A technicality

- Suppose you estimate $\Pr(\beta_j < 0) = 0.98$, $\Pr(\beta_j > 0) = 0.01$, $\Pr(\beta_j = 0) = 0.01$.
- Should you declare an fdr of 0.01 or 0.02?
- Maybe fsr makes more sense anyway?

Unification of one-group and two-group solutions

- A possible concern with one-group model: if we assume all β are non-zero, but really some are actually zero, might we not underestimate the fsr?

Unification of one-group and two-group solutions

- A possible concern with one-group model: if we assume all β are non-zero, but really some are actually zero, might we not underestimate the fsr?
- Consider our example, with $\Pr(\beta_j > 0) = 0.025$. If we actually allowed $\beta_j = 0$ then possibly all of this probability might actually land at $\beta_j = 0$.


Unification of one-group and two-group solutions

- A possible concern with one-group model: if we assume all β are non-zero, but really some are actually zero, might we not underestimate the fsr?
- Consider our example, with $\Pr(\beta_j > 0) = 0.025$. If we actually allowed $\beta_j = 0$ then possibly all of this probability might actually land at $\beta_j = 0$.
- I argue that, assuming symmetry of g near 0, this also provides an upper bound of how much of the $\Pr(\beta_j < 0) = 0.975$ might also move to 0.

Unification of one-group and two-group solutions

- A possible concern with one-group model: if we assume all β are non-zero, but really some are actually zero, might we not underestimate the fsr?
- Consider our example, with $\Pr(\beta_j > 0) = 0.025$. If we actually allowed $\beta_j = 0$ then possibly all of this probability might actually land at $\beta_j = 0$.
- I argue that, assuming symmetry of g near 0, this also provides an upper bound of how much of the $\Pr(\beta_j < 0) = 0.975$ might also move to 0.
- Therefore a more conservative estimate of the fsr might be 0.05 (or, more generally, double what you get allowing for point mass)

Example: BRCA data



Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .

Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .
- So for example we can easily compute fdrs for discoveries other than “non-zero” (eg compute $\Pr(|\beta_j| > 2|\hat{\beta}_j|)$).

Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .
- So for example we can easily compute fdrs for discoveries other than “non-zero” (eg compute $\Pr(|\beta_j| > 2|\hat{\beta}_j|)$).
- And use it to obtain point estimates and credible intervals for each β_j , taking account of information from all the other β_j .

Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .
- So for example we can easily compute fdrs for discoveries other than “non-zero” (eg compute $\Pr(|\beta_j| > 2|\hat{\beta}_j|)$).
- And use it to obtain point estimates and credible intervals for each β_j , taking account of information from all the other β_j .
- Because $f(\beta)$ is unimodal, the point estimates will tend to be “shrunk” towards the overall mean (0).


Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .
- So for example we can easily compute fdrs for discoveries other than “non-zero” (eg compute $\Pr(|\beta_j| > 2|\hat{\beta}_j|)$).
- And use it to obtain point estimates and credible intervals for each β_j , taking account of information from all the other β_j .
- Because $f(\beta)$ is unimodal, the point estimates will tend to be “shrunk” towards the overall mean (0).
- Because $f(\beta)$ is estimated from the data, the amount of shrinkage is adaptive to the data. And because of the role of s_j , the amount of shrinkage adapts to the information on each gene.


Estimation and Shrinkage

- Besides allowing one to estimate fdr and fsr , this approach also provides a full posterior distribution for each β_j .
- So for example we can easily compute fdrs for discoveries other than “non-zero” (eg compute $\Pr(|\beta_j| > 2|\hat{\beta}_j|)$).
- And use it to obtain point estimates and credible intervals for each β_j , taking account of information from all the other β_j .
- Because $f(\beta)$ is unimodal, the point estimates will tend to be “shrunk” towards the overall mean (0).
- Because $f(\beta)$ is estimated from the data, the amount of shrinkage is adaptive to the data. And because of the role of s_j , the amount of shrinkage adapts to the information on each gene.
- So we call the approach “Adaptive Shrinkage” (ASH).


Example: ASH applied to mouse data




Example: ASH applied to mouse data




Example: ASH applied to mouse data



Shrinkage is adaptive to information




Shrinkage is adaptive to information




Shrinkage is adaptive to information

##	gene	lv1	lv2	rv1	rv2	pval	zdat.ash\$localfdr
## 19422	Mgat5b	7	10	320	452	0.03795	0.74069
## 20432	Sec63	1042	1034	5496	6649	0.04908	0.06952

Recall FDR problem 1: q values increased by low count genes



ASH q values more robust to inclusion of low count genes



Summary: FDR via Adaptive Shrinkage

- Both provide a rational approach to identifying “significant” findings.

Summary: FDR via Adaptive Shrinkage

- Both provide a rational approach to identifying “significant” findings.
- Both are generic and modular: once you have the summary data, you can forget where they came from.

Summary: FDR via Adaptive Shrinkage

- Both provide a rational approach to identifying “significant” findings.
- Both are generic and modular: once you have the summary data, you can forget where they came from.
- But by using two numbers ($\hat{\beta}, s$) instead of one (p values) precision of different measurements can be better accounted for.

Summary: FDR via Adaptive Shrinkage

- Both provide a rational approach to identifying “significant” findings.
- Both are generic and modular: once you have the summary data, you can forget where they came from.
- But by using two numbers ($\hat{\beta}, s$) instead of one (p values) precision of different measurements can be better accounted for.
- ASH borrows information for estimation, as well as testing.

Other Applications

- Widely applicable: perhaps anywhere (?) where shrinkage is appropriate, requiring only an estimated effect size and standard error for each object.

Other Applications

- Widely applicable: perhaps anywhere (?) where shrinkage is appropriate, requiring only an estimated effect size and standard error for each object.
- Could also use effect size estimate and p value for each variable, by converting to effect size estimate and (pseudo-) standard error.

Other Applications

- Widely applicable: perhaps anywhere (?) where shrinkage is appropriate, requiring only an estimated effect size and standard error for each object.
- Could also use effect size estimate and p value for each variable, by converting to effect size estimate and (pseudo-) standard error.
- Currently applying it to wavelet shrinkage applications.

Guarantees?

- “I think you have some nice ideas. How will you convince people to use them?” (C Morris)

Next steps?

- Extend to allow $g(\cdot; \pi)$ to depend on covariates X .

Next steps?

- Extend to allow $g(\cdot; \pi)$ to depend on covariates X .
- Extend to allow for correlations in the measured $\hat{\beta}_j$.

Thanks

- to the several postdoctoral researchers and students who have worked with me on related topics.

Thanks

- to the several postdoctoral researchers and students who have worked with me on related topics.
- Especially Mengyin Lu who coded the VB algorithm.

Reproducible research

- This document is produced with **knitr**, **Rstudio** and **Pandoc**.

Reproducible research

- This document is produced with **knitr**, **Rstudio** and **Pandoc**.
- For more details see my `stephens999/ash` repository at <http://www.github.com/stephens999/ash>

Reproducible research

- This document is produced with **knitr**, **Rstudio** and **Pandoc**.
- For more details see my `stephens999/ash` repository at <http://www.github.com/stephens999/ash>
- Website: <http://stephenslab.uchicago.edu>

Pandoc Command used

```
pandoc -s -S -i --template=my.beamer -t beamer -V  
theme:CambridgeUS -V colortheme:beaver ilike-slides.md -o  
ilike-slides.pdf
```

(alternative to produce html slides; but figures would need reworking)

```
pandoc -s -S -i -t dzslides --mathjax NSmeet2013.md -o  
NSmeet2013.html
```

Here is my session info:

```
print(sessionInfo(), locale = FALSE)  
  
## R version 3.0.2 (2013-09-25)  
## Platform: x86_64-apple-darwin10.8.0 (64-bit)  
##  
## attached base packages:  
## [1] splines      parallel     stats        graphics     grDevices    utils  
## [8] methods     base
```

Some odd things in the data




Figure : plot of chunk unnamed-chunk-43

```
##           gene lv1 lv2 rv1   rv2 genelength
## 17711  Napsa    0  1   7   779           1470
```